

Evaluating and Adapting Small Vision-Language Models for Document Understanding

Ali Rasekh

Abstract

We evaluate three small open vision-language models (VLMs) on a frozen DocVQA [Mathew et al., 2021] design with development-selected InternVL3-1B [Zhu et al., 2025, OpenGVLab, 2025], then compare a bounded language-attention LoRA [Hu et al., 2022] adaptation under a sealed holdout. Inference uses a custom PyTorch runner inspired by VLMEvalKit [Duan et al., 2024] DocVQA conventions, not a full VLMEvalKit benchmark execution. On development primary ANLS, pretrained InternVL3-1B outperforms SmolVLM-256M/500M [Marafioti et al., 2025] and is selected as the deployment candidate (step 0). Corrected LoRA v2 (update 200) is the fixed adaptation comparison. On holdout ($n=504$), base ANLS is approximately 0.7927 versus adapted 0.7993 (paired $\Delta=+0.006639$). A UCSF-cluster bootstrap 95% interval of approximately $[-0.00530, +0.01905]$ includes zero: under this analysis, no clear improvement is established, and we make no equivalence claim. Post-holdout supplementary Pascal development analyses (Appendix D) further probe visual-tile budgets, three-seed LoRA variability, and dependence on the correct document image; they do not change the sealed selection or official holdout result.

1 Introduction and research questions

Document visual question answering requires reading text, layout, and tables from page images [Mathew et al., 2021]. Small open VLMs enable reproducible cluster experiments, but published numbers often mix evaluation kits, preprocessing budgets, and metric variants. This report answers four questions under a frozen protocol:

1. How do three pretrained configured systems compare on a fixed DocVQA development split under a shared primary metric?
2. Does bounded language-attention LoRA of the selected InternVL system improve development primary ANLS enough to replace step 0?
3. On a once-evaluated holdout, how do the four systems score?
4. Does a paired UCSF-cluster bootstrap support a clear adapted-minus-base gain?

Comparisons are *configured-system* comparisons (weights, native preprocess, decoding, prompts), not isolated architecture effects.

Model selection (eligibility, not a survey). We restricted candidates to small open instruct VLMs compatible with our pinned PyTorch/Transformers stack and a roughly sub-billion-parameter operating budget. SmolVLM-256M and SmolVLM-500M [Marafioti et al., 2025] provide two sizes within one Idefics3-style family with documented efficient multimodal recipes. InternVL3-1B [Zhu et al., 2025, OpenGVLab, 2025] adds a distinct InternVL family with an official dynamic document-oriented visual preprocess. Together they cover family diversity and documented document/OCR-relevant capability without an exhaustive model survey.

From development observations to a bounded adaptation hypothesis. On development runs we observed recurring error modes: field/table-cell selection mistakes (wrong column or nearby cell), entity/text confusions among co-occurring names or phrases, and answer-format sensitivity (especially trailing punctuation for SmolVLM). These observations motivate a *hypothesis*, not a guarantee: language-attention LoRA may adjust task-conditioned answer selection and short-phrase formatting while leaving vision features fixed. Freezing the vision backbone and connector deliberately limits the method’s ability to fix visual-detail or extraction failures that require better visual encoding. Under our compute budget, adapting only language attention on a 2000-example subset for 200 updates was practical and informative: it tests whether a cheap, localized adaptation moves primary development ANLS enough to displace the pretrained baseline.

2 Task, data roles, and metrics

Why DocVQA. DocVQA provides short extractive answers over real document pages with multi-reference answers [Mathew et al., 2021].

One dataset, several roles. All subsets come from HuggingFaceM4/DocumentVQA revision `a44195f8...` (validation for smoke/dev/holdout; train shard for the adaptation subset). Table 1 lists roles; counts were verified against frozen manifests. This is **not** a collection of independent benchmarks.

Table 1: DocVQA roles used in this project (manifest-verified).

Role	Questions	UCSF groups	Notes
Smoke	8	7	Pipeline check
Development	208	45	Selection / comparison
Train (prepared)	2000	298	Adaptation pool; 1600 unique examples seen
Robustness (dev)	54	16	Synthetic stress subset of development
Holdout v2	504	109	Sealed final evaluation

Split construction (accurate history). Initial frozen splits were sampled with `doc_id` grouping. A later source-group audit found UCSF-level overlap between development and holdout v1. **Development membership was retained** after that audit; **holdout was rebuilt** as `holdout_v2` to remove UCSF overlap with smoke/development [Mathew et al., 2021]. We do not claim both splits were originally sampled under the corrected UCSF grouping scheme. Image integrity (PNG hashes) was audited for the frozen sets.

Relation to VLMEvalKit. VLMEvalKit [Duan et al., 2024] documents a DocVQA_VAL route and ANLS-style utilities. Our evaluation is a custom greedy runner with frozen manifests and a project primary metric. We do **not** claim a full VLMEvalKit leaderboard run. Secondary diagnostics (VLMEvalKit-compatible ANLS; strip ≤ 1 trailing period after normalize) are not used for selection.

Primary metric (self-contained). Primary score is project metric `anls_normalized_strict_v2` (normalized strict ANLS), **not** claimed byte-identical to official Pix2Struct [Lee et al., 2023] or

VLMEvalKit scorers. Normalization: lowercase, strip ends, collapse internal whitespace to single spaces. Let NL be Levenshtein distance divided by $\max(|\hat{y}|, |y|)$ on the normalized prediction and reference strings. Per-reference score is $1 - \text{NL}$ if $\text{NL} < 0.5$, else 0. Example score is the maximum over valid nonempty references. Primary EM uses the same normalization and requires equality to some reference (no punctuation stripping). Failures and missing predictions count as zero in requested-set means over the full question set.

3 Systems and visual preprocessing

Table 2: Pretrained configured systems. Parameter counts are measured unique parameters.

System	Unique params	Notes
SmolVLM-256M-Instruct	256,484,928	Idefics3-style
SmolVLM-500M-Instruct	507,482,304	Idefics3-style
InternVL3-1B	938,193,024	Dynamic 448px tiles

Native visual preprocessing differs across families and yields unequal visual-token budgets. SmolVLM uses processor defaults; InternVL3-1B uses the official dynamic preprocess (448px tiles, `max_num` = 12, thumbnail enabled) from the model card [OpenGVLab, 2025].

InternVL3-1B backbone names. The official InternVL3-1B model card [OpenGVLab, 2025] identifies vision backbone `InternViT-300M-448px-V2.5` and language backbone `Qwen2.5-0.5B`. Runtime Transformers architecture notes report class `Qwen2ForCausalLM`; that class name is an implementation detail, not a reason to name the backbone simply “Qwen2”. Vision runtime module is `InternVisionModel` within `InternVLChatModel`; connector uses downsample ratio 0.5 (pixel-unshuffle style).

4 Model profiles, qualitative gaps, and measurement plan

Model profiles (documented sources). SmolVLM-256M/500M [Marafioti et al., 2025] are compact Idefics3-style instruct VLMs. The public Hugging Face model card for SmolVLM-256M-Instruct states that training data come from **The Cauldron** and **Docmatix**, with emphasis on document understanding ($\approx 25\%$) and image captioning ($\approx 18\%$), plus broader multimodal instruction mixtures; a full exhaustive dump of mixture members is not provided on the card. InternVL3-1B [Zhu et al., 2025, OpenGVLab, 2025] pairs `InternViT-300M-448px-V2.5` with a `Qwen2.5-0.5B`-initialized language component and documents **Native Multimodal Pre-Training** that interleaves multimodal sequences with large-scale text corpora (card/paper-level description). The InternVL3 card does **not** publish a complete dataset inventory sufficient to prove or disprove DocVQA contamination. We therefore report the disclosed mixtures/methods above and still treat residual pretrained exposure as an open risk (Limitations), not a settled fact.

Qualitative knowledge-gap sample (development; descriptive only). From InternVL3-1B development predictions we drew a seeded sample of $n=20$ non-perfect examples (48 non-perfect among 208; selection rule in `results/internvl_error_sample.v1.json`). This sample is **qualitative and not a quantitative error taxonomy** for the full split. Recurring observed tags in the inspected sample include text/checkbox reading mistakes, nearby field or table-cell

confusions, and formatting/entity mismatches between prediction and references. These observations motivated the *working* hypothesis that language-attention LoRA might help short-phrase formatting and answer selection while leaving vision features fixed—a post-hoc interpretation of development errors, **not** a pre-registered experimental hypothesis.

Expected outcomes and measurement criteria (stated retrospectively). Under the frozen protocol we expected: (i) InternVL to lead SmolVLM on development primary ANLS given its stronger document preprocess; (ii) bounded LoRA either to improve development ANLS enough to displace step 0 or to show only small movements; (iii) holdout to report absolute configured-system scores with a paired InternVL adapted–base interval that may include zero. Measurement criteria remained primary `anls_normalized_strict_v2` on frozen requested sets, with EM secondary and formatting diagnostics non-selective. Improvement plan if revisiting adaptation: larger multi-seed budgets, matched visual-token settings across families, and diagnostics that isolate visual grounding beyond whole-page replacement (Appendix D).

Software stack and edge extrapolation. The tested stack is PyTorch 2.5.1+cu118 with Transformers 4.48.3 and PEFT 0.14.0 on server GPUs (Neumann A100 / Pascal V100 allocations as recorded per experiment). Latency and peak-memory figures characterize those server allocations only; they do **not** transfer to edge or mobile devices without re-measurement under the target runtime. Generic CPU-only `pip install` environments are insufficient to reproduce GPU timings or bf16 training behavior.

5 Development comparison and selection

Table 3: Development primary metrics ($n=208$). Diagnostic EM strips ≤ 1 trailing period after normalize (secondary).

System	ANLS	EM	Diag. EM
SmolVLM-256M	0.5568	0.2596	0.4952
SmolVLM-500M	0.6095	0.0433	0.6298
InternVL3-1B (step 0)	0.8324	0.7692	0.7692

Table 3 summarizes pretrained development scores on Neumann (InternVL base ANLS 0.8324). Pascal supplementary matched-base ANLS 0.8276 is a separate recorded run; see §11. Under the pre-declared rule (primary development ANLS; earlier step on ties), InternVL step 0 is the deployment candidate. Corrected LoRA v2 scored ≈ 0.8263 (update 100) and ≈ 0.8290 (update 200), both below step 0. Best corrected trained checkpoint is update 200 (adaptation comparison only). **Selection protocol:** freeze deployment = step 0 and adaptation comparison = LoRA v2 u200 from development; holdout rankings are reported but do not change that selection.

6 Adaptation recipe




Figure 1: InternVL3-1B with LoRA on language-model attention. Backbone names follow the official model card [OpenGVLab, 2025].




Figure 2: Data preparation → LoRA training → development selection → robustness subset → sealed holdout.

LoRA [Hu et al., 2022] targets `language_model` attention `q_proj/k_proj/v_proj/o_proj` ($r=16$, $\alpha=32$, dropout 0.05), freezing vision and connector. Adapted total unique parameters including adapters: 940,355,712; trainable: 2,162,688.

Supervised labels (corrected v2). Historical v1 training duplicated the assistant end-of-turn token (2000/2000 audited examples). v1 is historical/superseded; we do not claim it caused every score difference, and we do not compare v1/v2 loss magnitudes as the same objective. Corrected v2 uses a single assistant EOS; labels supervise the assistant answer span only (prompt tokens masked).

Bounded training (actual config). Target answer = first nonempty source-listed answer. Optimizer AdamW (lr 5×10^{-5} , weight decay 0.01); linear warmup schedule over 10 updates then decay across 200 updates; grad clip 1.0; microbatch 1; grad accum 8; bf16 when supported; seed 42; max sequence length 8192. From the prepared 2000-example subset, training presented 1600 sample presentations (1600 unique QIDs under shuffle+wrap). Figure 3: first-20 mean ≈ 0.2561 , last-20 mean ≈ 0.2143 (full precision retained in artifacts).




Figure 3: Corrected LoRA v2 per-update training loss with numeric axis ticks.

7 Development robustness (limited)

On a fixed development subset ($n=54$ / 16 UCSF groups) we apply three image conditions before native preprocess: *clean*; *half-detail* (BICUBIC downsample by $\text{floor}(w/2)$, $\text{floor}(h/2)$, then restore exact original dimensions); *JPEG q40* (quality 40, Pillow `subsampling=2` / 4:2:0). Preserving dimensions keeps InternVL dynamic tile counts comparable across conditions so the stress is content degradation, not a different tiling budget. Figure 4 shows small differences; this is not a corruption benchmark. Phase 3B v1 robustness figures remain historical only.




Figure 4: Development robustness primary ANLS (dot/line; y-axis truncated to [0.70, 0.80]).

8 Holdout results

Protocol freeze revision `a9b0695` . . . ; four systems once on holdout_v2 ($n=504$; job 84636; A100-PCIE-40GB). Latency is CUDA-synchronized *generation-only* time (warmup excluded), not end-to-end wall-clock. Table 4 and Figure 5 report primary metrics; Appendix B adds p95 latency and peak allocated memory.

Table 4: Holdout primary metrics ($n=504$). Values are fractions in $[0, 1]$.

System	ANLS	EM	Med. gen. (s)
SmolVLM-256M	0.5753	0.3234	0.214
SmolVLM-500M	0.5997	0.0694	0.244
InternVL3-1B (base)	0.7927	0.7282	0.482
InternVL + LoRA v2 u200	0.7993	0.7401	0.538




Figure 5: Holdout primary ANLS for the four configured systems.

Secondary VLMEvalKit-style ANLS preserves the InternVL base/adapted ordering. SmolVLM-500M shows formatting sensitivity: primary EM ≈ 0.0694 vs diagnostic EM ≈ 0.6052 .

Paired adaptation analysis. Point estimate adapted–base ANLS = +0.006639 (about +0.66 percentage points if mapped to a 0–100 scale; we report fractions). UCSF-cluster bootstrap ($B=2000$, seed 42) yields 95% CI approximately $[-0.00530, +0.01905]$ (Figure 6). The interval includes zero: no clear improvement is established; we do not claim equivalence.




Figure 6: Paired holdout ANLS difference with UCSF-cluster 95% CI.

9 Illustrative input/output examples

Examples are deterministically selected (gallery manifest) and illustrative, not performance-representative. Crops are marked; full-page context remains available under `data/images/docvqa_holdout_v2/`

Shared table-column error (qid 307). Duration for “Replace Flex Tester” is 6 months; both systems predict 1.0 (PRIORITY cell). Figure 7 shows a labeled crop of the relevant rows.

Adaptation improvement (qid 18564). Invoice addressee question: base predicts “Dr. Samuel Deadwyler” (body/description name); adapted predicts “Dr. J. H. Robinson” (recipient box). Figure 8 pairs a marked full page with the addressee crop.

Adaptation regression (qid 792). Letter mentioning a free sample of SAMSON: base answers “SAMSON”; adapted answers “free sample”. Figure 9 pairs marked full page and sentence crop.

57036347

Appendix IIB
Methods Development Two-Year Project Plan

PROJECTS TO BE DELEGATED OR DROPPED	R&D CUSTOMER	PRIORITY	DURATION MONTHS	PROJECT LEADER
1. Replace Flex Tester	METH DEV	1.0	6	FY
2. Water on Tobacco	METH DEV	2.0	4	TVB
3. Develop Filter Hardness Method	AM	2.1	4	FY
4. RAA Starch on Tobacco	METH DEV	3.0	4	NMD3
5. R&D Simile vs ISO Method for H2O/Nicotine	KRW	4.0	3	NMD2
6. RAA Inorganic Ions	METH DEV	4.0	8	NMD3
7. Develop Filter Hardness Method	KRK	6.0	4	FY
8. Develop Hot Collapse Hardness Method	KRK	7.0	4	NMD2
9. Water Activity on Snuff	RRB	8.0	6	TVB
10. Insulin Tab. Firmness/Fill Value	AM	10.0	6	FY
11. RAA Organic Tobacco Acids	RRB	12.0	12	NMD3
12. HPLC Methoprene Method for Tobacco	RRB	13.0	4	TVB

* NMD1 = STH Replacement
* NMD2 = Robotics/Analytical Chemist
* NMD3 = VLO Replacement

Source: <https://www.industrydocuments.ucsf.edu/docs/hrcj0037>

Full page (crop box marked)

Methods Development Two-Year Project Plan

PROJECTS TO BE DELEGATED OR DROPPED	R&D CUSTOMER	PRIORITY	DURATION MONTHS	PROJECT LEADER
1. Replace Flex Tester	METH DEV	1.0	6	FY
2. Water on Tobacco	METH DEV	2.0	4	TVB
3. Develop Filter Hardness Method	AM	2.1	4	FY
4. RAA Starch on Tobacco	METH DEV	3.0	4	NMD3

Evidence crop

Figure 7: Illustrative shared error (qid 307): PRIORITY vs DURATION.

THE BOWMAN GRAY SCHOOL OF MEDICINE
OF WAKE FOREST UNIVERSITY
300 South Hawthorne Road, Winston-Salem, North Carolina 27103
Office of the Controller
(919) 746-4445

WAKE FOREST UNIVERSITY

Dr. J. H. Robinson
Bowman Gray Technical Center
R. J. Reynolds Tobacco Co.
Winston-Salem, NC 27102

Invoice No. G-1704
Date 10/9/91

CASHIER
Credit 2-144-831-1140-751
Credit
Credit

INVOICE
REFER TO INVOICE NUMBER WHEN PAYING

Date	Ref.	Description	Total
		Contract #91-771-014 Support of "Investigation of the Effects of Nicotine in the Mammalian Brain" Dr. Samuel Deadwyler Quarterly Billing for October 1, 1991 - December 31, 1991	\$17,019.25

50820 2277

Full page (crop box marked)

CROP addressee (Dr. J. H. Robinson)

WAKE FOREST UNIVERSITY

Dr. J. H. Robinson
Bowman Gray Technical Center
R. J. Reynolds Tobacco Co.
Winston-Salem, NC 27102

Addressee crop

Figure 8: Illustrative adaptation improvement (qid 18564).

The Coca-Cola Company
ATLANTA, GEORGIA

July 22, 1974

ADDRESS REPLY TO
P. O. BOX 1734
ATLANTA, GA. 30301
404-887-2121

Mr. Robert White
Center for Science in the
Public Interest
1779 Church Street, N.W.
Washington, D. C. 20036

Dear Mr. White:

Your letter of June 8, 1974, addressed to The Coca-Cola Company Foods Division has been referred to me. Please accept my apologies for this tardy response as your letter has just been brought to my attention.

Enclosed please find a free sample, as you requested, of our orange flavored nutritional beverage powder, SAMSON. Most of the information pertaining to the nutritional qualities of this product appears on the label, however, a short fact sheet is enclosed giving a synopsis of our efforts developing nutritional beverages from 1966 to the present.

On behalf of The Coca-Cola Company I wish to express appreciation for your interest in SAMSON. If further information or samples are needed please feel free to contact me.

Respectfully yours,
G. L. K. Hunter
G. L. K. Hunter

GLKH:nh

cc: Dr. D. S. Bhatia

Source: <https://www.industrydocuments.ucsf.edu/docs/hfpj0226>

Full page (crop box marked)

CROP letterhead of HUNTER

Dear Mr. White:

Your letter of June 8, 1974, addressed to The Coca-Cola Company Foods Division has been referred to me. Please accept my apologies for this tardy response as your letter has just been brought to my attention.

Enclosed please find a free sample, as you requested, of our orange flavored nutritional beverage powder, SAMSON. Most of the information pertaining to the nutritional qualities of this product appears on the label, however, a short fact sheet is enclosed giving a synopsis of our

Sentence crop

Figure 9: Illustrative adaptation regression (qid 792).

10 Limitations and future work

- Sampled DocVQA evaluation, not official full-benchmark leaderboard parity.
- Unequal native visual-token budgets across families.
- Possible pretrained exposure to DocVQA-related data is not ruled out from public cards alone.
- Original selection used a single seed and bounded LoRA budget; historical v1 double-EOT training superseded.
- Phase 3B robustness evidence is limited to the fixed synthetic development subset (clean / half-detail / JPEG q40); it is distinct from the later visual-input-dependence study.
- Bootstrap CIs describe this holdout resampling design only; they are not seed SDs.
- Supplementary Pascal runs used V100 BF16 usability (not native $CC \geq 8$) and, for P3, a documented cross-node GPU SKU change (Appendix D).

11 Supplementary Pascal development analyses

After sealing holdout selection, we ran *post-holdout* development analyses on Pascal under companion ownership (`companion/pascal/`). They did **not** change the deployment candidate (InternVL step 0) or official holdout scores. Details and figures are in Appendix D; compact evidence is under `companion/pascal/results/`.

Visual-budget ablation (P1). On the frozen InternVL3-1B **base** (no LoRA), reducing only `max_num` from 12→4→1 on development ($n=208$) lowers primary ANLS from approximately 0.8276 to 0.7588 to 0.5595, with corresponding reductions in actual tile usage and request latency on the recorded V100 run. This is a within-runtime quality/cost trade-off under matched decoding, not a claim about other GPUs. The Pascal matched-base figure 0.8276 is a *separate recorded run* from the earlier Neumann development InternVL base score 0.8324 in Table 3; documented differences include host/GPU (Pascal V100 vs Neumann A100), allocation, and job identity. We do **not** assert a single verified cause for that score gap; Pascal tables use the Pascal matched control only.

Three-seed fixed-budget LoRA (P3). Repeating corrected v2 settings for seeds {42, 43, 44} to update 200 yields development ANLS {0.8193, 0.8178, 0.8317} versus a matched fresh base 0.8276. Arithmetic mean 0.8229; **sample SD across three seeds** 0.0076 (ddof=1; *not* a confidence interval). Two seeds fall below base and one is slightly above: under this fixed configuration and budget we observe **no consistent improvement** over the matched base. This does **not** claim that LoRA generally fails. Each completed run saw 1600 unique QIDs from the 2000-example pool (80% coverage); sample order/membership differs by seed. These Pascal adapters are separate from the official Neumann v2 u200 adaptation comparison.

Visual-input dependence. On the pinned InternVL3-1B base only, replacing the correct page with a white or unrelated development image reduced development ANLS by approximately **78 percentage points**, indicating strong dependence on the correct visual input. That summary is not an exact decomposition of performance, not proof of fine-grained grounding, and not evidence for or against contamination. White residual ANLS ≈ 0.042 ; unrelated ≈ 0.047

(scored against the *original* question references). Prediction agreement with the clean prediction is a separate diagnostic from correctness; agreement uses only pairs where both records succeed.

12 Conclusion

Under the frozen protocol, InternVL3-1B step 0 remains the development-selected deployment candidate; corrected LoRA v2 update 200 is the fixed adaptation comparison. Holdout shows a small positive ANLS point estimate whose UCSF-cluster 95% CI includes zero. Configured-system gaps between SmolVLM and InternVL are large; adaptation gains under the original budget are not clearly established. Supplementary Pascal analyses reinforce a strong visual-input dependence for the selected base and show no consistent three-seed LoRA gain under the fixed corrected recipe, without revising sealed selection.

A Reproducibility pins

Item	Pin
Python / torch / transformers / peft	3.11 / 2.5.1+cu118 / 4.48.3 / 0.14.0
SmolVLM-256M	HuggingFaceTB/SmolVLM-256M-Instruct @ 7e3e67ed...
SmolVLM-500M	HuggingFaceTB/SmolVLM-500M-Instruct @ a7da5b98...
InternVL3-1B	OpenGVLab/InternVL3-1B @ 4415a3b8...
Vision / language (card)	InternViT-300M-448px-V2.5 / Qwen2.5-0.5B
DocVQA HF mirror	HuggingFaceM4/DocumentVQA @ a44195f8...
VLMEvalKit	54a063c5... (route/metric reference only)
reference rev	
Eval protocol freeze	git a9b0695...
LoRA v2 u200	3c2c34cd...c344180a
adapter SHA256	
Holdout job	84636

Commands: `reports/README.md`, `docs/report/SUBMISSION_GUIDE.md`. Official LoRA v2 u200 adapter weights are tracked under `artifacts/internvl3_1b_lora_v2_u200/` (SHA256 in `artifacts/internvl3_1b_lora_v2_u200/`). base weights and DocVQA images must still be downloaded separately.

B Holdout latency and memory

Table 5: Generation-only latency and peak allocated GPU memory (holdout; from frozen summaries).

System	Med. (s)	p95 (s)	Peak alloc. (GiB)
SmolVLM-256M	0.214	0.410	2.78
SmolVLM-500M	0.244	0.427	3.26
InternVL3-1B (base)	0.482	0.647	3.42
InternVL + LoRA v2 u200	0.538	0.799	3.43

C Historical v1 note

v1 LoRA used duplicated supervised EOT labels; superseded by v2. Do not compare v1/v2 losses as the same objective. Chronology: docs/progress/phase03c.md.

D Supplementary Pascal evidence

All tables below use the development requested set ($n=208$) unless noted. Failures/missing count as zero; generation-cap hits remain in requested-set means. These analyses are labeled **post-holdout supplementary development work**.

A. Visual-budget ablation (P1)

Pinned InternVL3-1B base; conditions `max_num` $\in \{12, 4, 1\}$; greedy; `max_new_tokens` = 64; Pascal V100 (BF16 usability). Job 117225.

Table 6: P1 development primary metrics by tile-budget cap (requested-set means).

<code>max_num</code>	ANLS	EM	Mean tiles (observed)
12	0.8276	0.7644	12.58
4	0.7588	0.6538	4.62
1	0.5595	0.4183	1.00




Figure 10: P1: development ANLS versus observed mean tile usage under `max_num` caps (InternVL3-1B base).

B. Three-seed corrected LoRA (P3)

Settings match corrected v2 (seeds 42/43/44; 200 updates; accum 8). Fresh-base and adapter evaluations share the development manifest denominator. Figure 11 marks the common fresh base; Figure 12 shows raw loss (faint) with trailing 20-update means (samples may differ by seed).

Table 7: P3 development ANLS/EM (requested set). Sample SD is across three seed ANLS values (ddof=1), not a CI.

System	ANLS	EM	Δ ANLS vs base	$\uparrow/\downarrow/=$
Fresh base	0.8276	0.7644	—	—
Seed 42	0.8193	0.7644	−0.0084	5/6/197
Seed 43	0.8178	0.7596	−0.0098	8/7/193
Seed 44	0.8317	0.7740	+0.0040	6/3/199
3-seed mean / sample SD	0.8229	—	—	SD 0.0076




Figure 11: P3 development ANLS for three Pascal seeds versus fresh base (dashed) and three-seed mean (dotted).




Figure 12: P3 training loss: low-opacity raw traces and trailing 20-update means. Training examples can differ across seeds.

Coverage and provenance. Each completed run: 200 optimizer updates, 1600 presentations, 1600 unique QIDs (80% of the 2000-pool). Protocol deviation: seeds 42/43 trained on pascal-node10 (V100-SXM3); completed seed 44 and all four evaluations on pascal-node09 (V100-PCIE). Interrupted 164-update seed 44 attempt preserved and excluded. Official Neumann v2 u200 remains a separate historical reference.

C. Visual-input dependence

Pinned InternVL3-1B **base** only (not a P3 adapter); conditions clean / white (exact original $W \times H$) / unrelated donor from the same 208-set. Job 117260; one allocation.

Table 8: Visual-input dependence on development ($n=208$ ok / 0 fail / 0 miss). Agreement denominator: pairs where both condition and clean have `status=ok` (here 208). Agreement $\neq$ correctness.

Condition	ANLS	EM	Δ ANLS vs clean	Agree vs clean
Clean	0.8276	0.7644	0	1.000
White	0.0421	0.0288	-0.7855	0.024
Unrelated	0.0469	0.0288	-0.7807	0.029




Figure 13: Visual-input dependence: requested-set ANLS and EM for clean, white, and unrelated conditions.

Interpretation. Replacing the correct page with a white or unrelated image reduced development ANLS by approximately 78 percentage points, indicating strong dependence on the correct visual input. Unrelated scores use the original question’s references (mismatch diagnostic; layout/tiles may change). White/clean tile histograms matched *in this run* (observation, not a general law). Residual correct answers do not establish memorization or contamination; the study does not prove fine-grained visual grounding.

Mapping deviation. Original protocol wording described ascending donor QID then residual seeded shuffle. Executed code: among equal minimum aspect-distance donors, seeded choice via `random.Random(2026).choice` (candidates ordered by QID). Frozen mapping retained; not regenerated.

QID 292
From which source the data is taken in this document?

Consumer Dynamics		
KOL		
		Index
• Retention of Franchisees	83.6%	(100)
• Rate of Switching Leases	10.3%	(87)
• Rate of Quitting Leases	6.1%	(73)
• Single Brand Users in the Franchise:	78.2%	(111)
• Share of Industry Switches Gained:	0.8%	(28)
• Share of the 25-25 Segment:	1.6%	(31)

Source: (2008) 100 cities, 10,000+ Data

Original input size: 2103×1604 px (PDF thumbnail is display-only).

USMM 1/95-6/95, 12-Month Data

Original input size: 2103×1604 px (PDF thumbnail is display-only).

Intentionally blank input; same original dimensions as the correct document.

WHO

TABLE IV-1-1. PROPERTY DISTRIBUTION OF CYCLOHEXANE POLYMER DEFERENCE SAMPLE BY GRAVIMETRY, AREA OF PEAKS, INFRARED 10, THERMAL ANALYSIS, NMR AND XRD, ANALYSIS										
ANALYSIS	Sample	Linear		Branched		Cyclic		Cross-linked		Residual
		wt. %	Area %	wt. %	Area %	wt. %	Area %	wt. %	Area %	
GRAVIMETRY	1	100	100	0	0	0	0	0	0	0
	2	100	100	0	0	0	0	0	0	0
	3	100	100	0	0	0	0	0	0	0
	4	100	100	0	0	0	0	0	0	0
	5	100	100	0	0	0	0	0	0	0
	6	100	100	0	0	0	0	0	0	0
	7	100	100	0	0	0	0	0	0	0
	8	100	100	0	0	0	0	0	0	0
	9	100	100	0	0	0	0	0	0	0
	10	100	100	0	0	0	0	0	0	0
AREA OF PEAKS	1	100	100	0	0	0	0	0	0	0
	2	100	100	0	0	0	0	0	0	0
	3	100	100	0	0	0	0	0	0	0
	4	100	100	0	0	0	0	0	0	0
	5	100	100	0	0	0	0	0	0	0
	6	100	100	0	0	0	0	0	0	0
	7	100	100	0	0	0	0	0	0	0
	8	100	100	0	0	0	0	0	0	0
	9	100	100	0	0	0	0	0	0	0
	10	100	100	0	0	0	0	0	0	0
INFRARED 10	1	100	100	0	0	0	0	0	0	0
	2	100	100	0	0	0	0	0	0	0
	3	100	100	0	0	0	0	0	0	0
	4	100	100	0	0	0	0	0	0	0
	5	100	100	0	0	0	0	0	0	0
	6	100	100	0	0	0	0	0	0	0
	7	100	100	0	0	0	0	0	0	0
	8	100	100	0	0	0	0	0	0	0
	9	100	100	0	0	0	0	0	0	0
	10	100	100	0	0	0	0	0	0	0
THERMAL ANALYSIS	1	100	100	0	0	0	0	0	0	0
	2	100	100	0	0	0	0	0	0	0
	3	100	100	0	0	0	0	0	0	0
	4	100	100	0	0	0	0	0	0	0
	5	100	100	0	0	0	0	0	0	0
	6	100	100	0	0	0	0	0	0	0
	7	100	100	0	0	0	0	0	0	0
	8	100	100	0	0	0	0	0	0	0
	9	100	100	0	0	0	0	0	0	0
	10	100	100	0	0	0	0	0	0	0
NMR	1	100	100	0	0	0	0	0	0	0
	2	100	100	0	0	0	0	0	0	0
	3	100	100	0	0	0	0	0	0	0
	4	100	100	0	0	0	0	0	0	0
	5	100	100	0	0	0	0	0	0	0
	6	100	100	0	0	0	0	0	0	0
	7	100	100	0	0					


Original input size: 2184×1682 px (PDF thumbnail is display-only).

Source: <https://www.industrydocuments.ucsf.edu/docs/snyc0227>

Figure 14: Illustrative visual-dependence example (qid 292; stacked clean / white / unrelated). Deterministic selection: clean EM=1, white EM=0; not a representative sample. The blank panel is an intentional white input (same native $W \times H$ as the correct page), not a missing figure.

QID 297
What is the percentage of single brand users in the franchise?

A. Correct document



The image shows a document titled "Consumer Dynamics KOOL" with a logo on the left. It contains a table with the following data:

	Index
• Retention of Franchisees:	85.6% (100)
• Rate of Switching Leases:	16.3% (87)
• Rate of Quitting Leases:	6.1% (72)
• Single Brand Users in the Franchise:	78.2% (111)
• Share of Industry Switchers Gained:	0.8% (26)
• Share of the 24-25 Segment:	1.6% (31)

Source: ISTAT 1994 data, 12-month data

Original input size: 2103×1604 px (PDF thumbnail is display-only).

Model prediction:

78.2%

B. Blank white input




Original input size: 2103×1604 px (PDF thumbnail is display-only).

Intentionally blank input; same original dimensions as the correct document.

Model prediction:

20

C. Unrelated document



The image shows a document with a table of statistical data. The table has multiple columns and rows, with some cells containing numerical values and others containing text. The table is titled "TABLE 1.1.1. SUMMARY OF STATISTICAL DATA BY REGIONAL AREA OF THE UNITED STATES, 1994-1995".

Original input size: 2184×1682 px (PDF thumbnail is display-only).


Model prediction:

2.4

Figure 15: Illustrative visual-dependence example (qid 297; same selection rule as Figure 14).

QID 5158
What is the Title of the document ?

A. Correct document



Original input size: 1666×2149 px (PDF thumbnail is display-only).

Model prediction:

Water Analysis

B. Blank white input




Original input size: 1666×2149 px (PDF thumbnail is display-only).

Intentionally blank input; same original dimensions as the correct document.

Model prediction:

SAS Institute Inc. 2008 Annual Report

C. Unrelated document



Original input size: 1782×2298 px (PDF thumbnail is display-only).


Model prediction:

management report

Figure 16: Illustrative visual-dependence example (qid 5158; same selection rule as Figure 14).

QID 5159
What is the factory name ?

A. Correct document




Original input size: 1666×2149 px (PDF thumbnail is display-only).

Model prediction:

Fort Morgan

B. Blank white input




Original input size: 1666×2149 px (PDF thumbnail is display-only).

Intentionally blank input; same original dimensions as the correct document.

Model prediction:

BASF

C. Unrelated document



Original input size: 1782×2298 px (PDF thumbnail is display-only).

Model prediction:

The Buettner Building

Figure 17: Illustrative visual-dependence example (qid 5159; same selection rule as Figure 14).

Reproduction pointers. Companion protocols: `docs/companion/pascal/`. Compact tables/figures: `companion/pascal/results/`. Primary inference/loading path remains the root project loader (`load_internvl + load_language_lora`); PEFT `base_model_name_or_path` is stale inherited metadata (see `docs/report/adapter_metadata_audit.md`).

References

- Haodong Duan, Junming Yang, Yuxuan Qiao, Xinyu Fang, Lin Chen, Yuan Liu, Xiaoyi Dong, Yuhang Zang, Pan Zhang, Jiaqi Wang, et al. VLMEvalKit: An open-source toolkit for evaluating large multi-modality models. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pages 11198–11201, 2024. doi: 10.1145/3664647.3685520.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=nZeVKeeFYf9>.
- Kenton Lee, Mandar Joshi, Iulia Raluca Turc, Hexiang Hu, Fangyu Liu, Julian Martin Eisenschlos, Urvashi Khandelwal, Peter Shaw, Ming-Wei Chang, and Kristina Toutanova. Pix2Struct: Screenshot parsing as pretraining for visual language understanding. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 18893–18912. PMLR, 2023. URL <https://proceedings.mlr.press/v202/lee23g.html>.
- Andrés Marafioti, Orr Zohar, Miquel Farré, Merve Noyan, Elie Bakouch, Pedro Cuenca, Cyril Zakka, Loubna Ben Allal, Anton Lozhkov, Nouamane Tazi, Vaibhav Srivastav, Joshua Lochner, Hugo Larcher, Mathieu Morlon, Lewis Tunstall, Leandro von Werra, and Thomas Wolf. SmolVLM: Redefining small and efficient multimodal models. *arXiv preprint arXiv:2504.05299*, 2025.
- Minesh Mathew, Dimosthenis Karatzas, and C. V. Jawahar. DocVQA: A dataset for VQA on document images. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 2200–2209, 2021.
- OpenGVLab. InternVL3-1B model card. <https://huggingface.co/OpenGVLab/InternVL3-1B>, 2025. Pinned revision 4415a3b810e636d11dfa86b0e9ba40bb00535aa8; vision InternViT-300M-448px-V2.5, language Qwen2.5-0.5B.
- Jinguo Zhu, Weiyun Wang, Zhe Chen, Zhaoyang Liu, Shenglong Ye, Lixin Gu, Hao Tian, Yuchen Duan, Weijie Su, Jie Shao, Zhangwei Gao, Erfei Cui, Xuehui Wang, Yue Cao, Yangzhou Liu, Xingguang Wei, Hongjie Zhang, Haomin Wang, Weiye Xu, Hao Li, Jiahao Wang, Nianchen Deng, Songze Li, Yinan He, Tan Jiang, Jiapeng Luo, Yi Wang, Conghui He, Botian Shi, Xingcheng Zhang, Wenqi Shao, Junjun He, Yingtong Xiong, Wenwen Qu, Peng Sun, Penglong Jiao, Han Lv, Lijun Wu, Kaipeng Zhang, Huipeng Deng, Jiaye Ge, Kai Chen, Limin Wang, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhao Wang. InternVL3: Exploring advanced training and test-time recipes for open-source multimodal models, 2025. URL <https://arxiv.org/abs/2504.10479>.